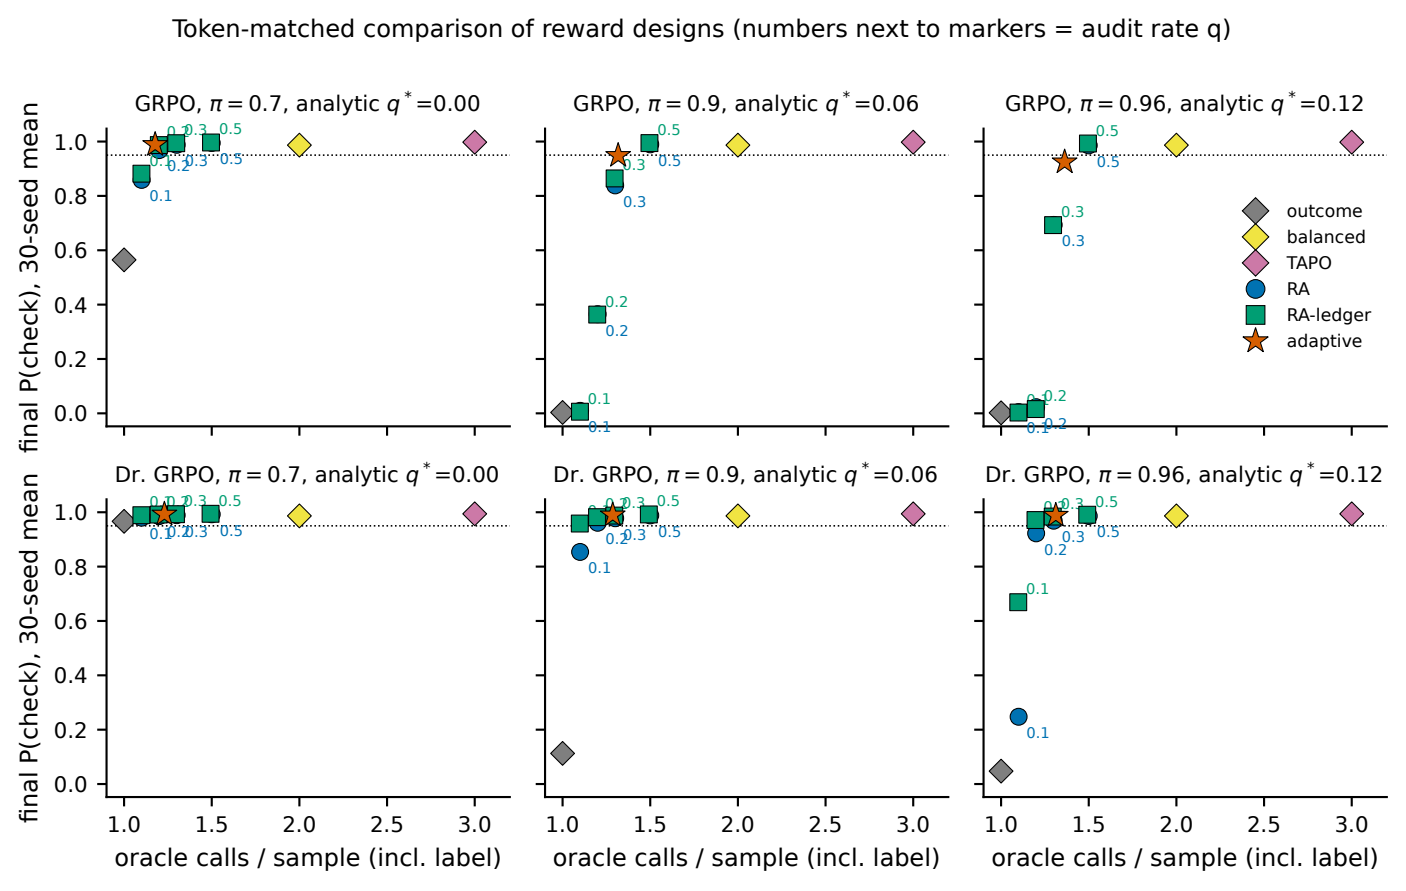
Token-matched comparison of reward designs (numbers next to markers = audit rate q)
GRPO, $\pi = 0.7$, analytic $q^* = 0.00$
GRPO, $\pi = 0.9$, analytic $q^* = 0.06$
GRPO, $\pi = 0.96$, analytic $q^* = 0.12$
Dr. GRPO, $\pi = 0.7$, analytic $q^* = 0.00$
Dr. GRPO, $\pi = 0.9$, analytic $q^* = 0.06$
Dr. GRPO, $\pi = 0.96$, analytic $q^* = 0.12$
final P(check), 30-seed mean
oracle calls / sample (incl. label)
outcome
balanced
TAPO
RA
RA-ledger
adaptive