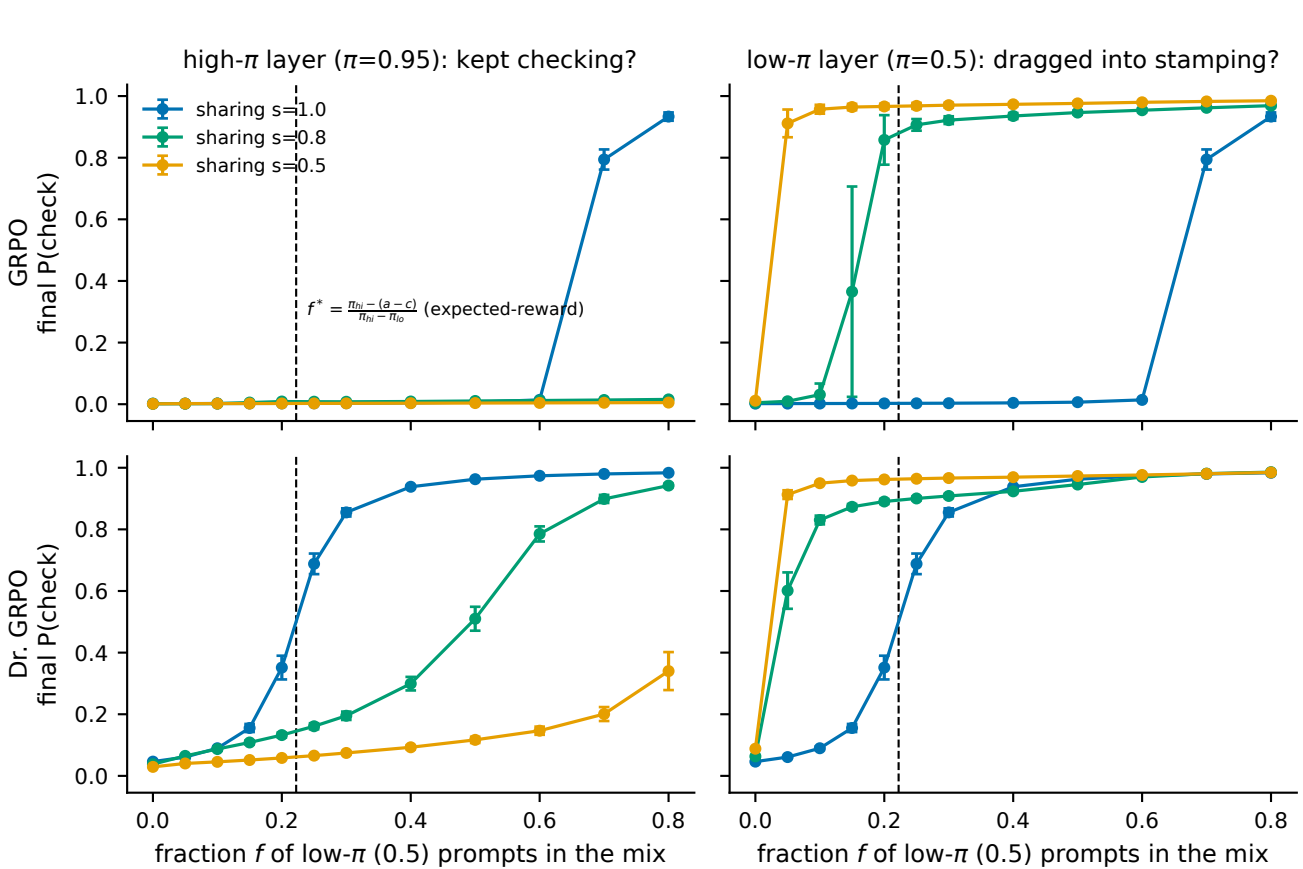
Data mixture as reward shaping (outcome-only reward, policy shared across layers)
high-π layer (π=0.95): kept checking?
low-π layer (π=0.5): dragged into stamping?
sharing s=1.0
sharing s=0.8
sharing s=0.5
$f^* = \frac{\pi_{hi} - (a-c)}{\pi_{hi} - \pi_{lo}}$ (expected-reward)
GRPO final P(check)
Dr. GRPO final P(check)
fraction f of low-π (0.5) prompts in the mix
fraction f of low-π (0.5) prompts in the mix